

Contents

1. Business Brief.....	2
2. Evidence ledger.....	3
3. Twelve-P diagnosis.....	4
4. The two priorities.....	6
5. Priority 1 — Profit.....	7
6. Priority 2 — Problem.....	9
7. Workflow backlog — the other ten Ps.....	11
8. Cadence.....	12
9. Execution receipt.....	12

XPMatch — Twelve-P Review

Run date: Sunday, Sep 21, 2026 · Framework: *Own or Be Owned* twelve Ps, via the own-run skill Stage: **pre-launch** — production software, zero commercial validation Branch: **idea** for everything commercial; **existing business** only for the software asset

Two disclosures before anything else. 1. **The twelve `own- specialist skills are not installed in this session.\*\* Only the own-run orchestrator was supplied. Every diagnosis below is therefore a \*\*provisional diagnostic built from references/p-map.md`, not the output of the specialist method for that P. Where a specialist would go deeper, the row says so.* 2. The source PDF (*Own or Be Owned*, 314 pages) was not attached. **The p-map's page citations could not be checked against it. No book anecdote is presented here as verified evidence. No official Owner Score is calculated — the book's questionnaire is not supplied, and inventing a score from a different rubric would misrepresent it. The 0–4 scale below is this suite's own evidence-based maturity scale, with unassessed and not yet applicable** as distinct states.**

1. Business Brief

Field	Value
Business	XPMatch, Inc. — taste-matched travel planning
Stage	Pre-launch. Production app; \$0 revenue, 0 customers, no billing integration
Customer	(A) self-planners taking 2+ leisure trips a year; (B) independent travel advisors, 0–3 years in the profession
Problem	Generic itineraries and 40 open tabs; advisors rebuild the same proposal for every client
Offer	Free · Trip Pass \$19/60d · Plus \$49/yr or \$7.99/mo · Concierge \$149–399 · Advisor seat \$49/mo
Currency / geography	USD, United States first
Owner path	Cash-flow positive by month 12; sellable by month 24 to a financial buyer
Constraints	\$25,000 founder cash [A]; 12 months with no owner draw [A]; ~50 founder hours/week combined [A]
Tools actually available	Railway, Postgres, Google Places, OpenRouter, Resend, GitHub. No Stripe, no CRM, no bookkeeping, no error monitoring, no scheduler
Current bottleneck	No mechanism to accept payment, compounded by zero first-party demand evidence
Team	Two founders [A]. No founder agreement, no vesting, no employees, no contractors
Last review	This document

2. Evidence ledger

ID	Evidence	Type	Source	Date / period
E1	Production application: onboarding, taste match, package builder, trip boards, maps, admin	Verified (first-party)	README.md, repository at 4320c89	Sep 21, 2026
E2	Cost per active user \$1.40–2.00/month after the place catalog, down from ~\$10	User-reported, first-party	docs/COGS.md	Sep 2026
E3	Competitor pricing: Layla \$49.99/yr · Wanderlog Pro \$39.99/yr · Triplt Pro \$49/yr · Travefy \$39–59/mo · Tern \$39/seat/mo	Verified (public pages, cited)	docs/COMPETITOR_LANDSCAPE.md, docs/REVENUE_MODEL.md	read Sep 2026
E4	78% of travel advisors charge a planning fee; fees \$100–1,500	Verified (published)	Travel Institute via Creo Proposals; AAA	2026
E5	37% of US travellers used AI to plan travel (2,001 adults); 43% of affluent travellers	Verified (published)	Allianz Mar–Apr 2026; Deloitte 2026	Spring 2026
E6	Fora: 15,000+ advisors, 97% new to the profession, \$1B valuation	Verified (published)	Skift, Fora newsroom	Jul 2026
E7	Expedia acquired Layla	Verified (published)	Expedia IR	Jul 31, 2026
E8	Travel subscription benchmarks: 4.1% median trial conversion; ~40% annual renewal; ~53% monthly first renewal	Verified (published)	RevenueCat State of Subscription Apps	2026
E9	Three-year projection: Y1 \$40,123 → Y3 \$427,320; Y3 SDE \$172,073; Y1 SDE –\$5,575	Derived calculation	2-financial-model.md, 2-financial-model.xlsx	Sep 21, 2026
E10	Postgres runs one instance, one volume, no backups; API keys pasted in chat remain unrotated	Verified (first-party)	docs/BETA_READINESS.md items 2 and 5	Sep 20, 2026
E11	No billing integration exists in the codebase	Verified (first-party)	Repository search; docs/BUSINESS_PLAN.md §6.5	Sep 21, 2026
E12	Zero paying customers, zero paid pilots, zero booked discovery calls, zero waitlist	Verified (absence, first-party)	docs/BUSINESS_PLAN.md §2 gate table	Sep 21, 2026
—	Community sizes, SEO head-term volumes, seasonality amplitude, support load per customer, name	Unknown — collection tasks, not zeros	—	—

ID	Evidence	Type	Source	Date / period
	clearance			

Not evidence: no customer interview, transcript, pre-order, letter of intent or testimonial exists for XPMatch. Nothing in E3–E8 is evidence that *this* product will be bought; it is evidence that the category transacts. Desk research is not customer validation.

3. Twelve-P diagnosis

Maturity: 0 = evidence confirms absent · 1 = ad hoc / owner-dependent · 2 = defined and trialled · 3 = repeatable, measured, delegated · 4 = sustained through exceptions and owner absence · **n/a** = not yet applicable at this stage · **?** = unassessed (not the same as 0).

#	P	Finding	Evidence	Maturity	Validation readiness	Impact	Key missing input
1	Pricing	Five offers priced against verified competitor anchors with per-offer margin math. Never spoken aloud to a buyer; no quote, discount or capacity history exists	E3, E4, E9	1	Medium	High	One real price conversation with a named buyer
2	Persona	Three segments documented with pain language, trigger events, watering holes and disqualifiers — but the pain language is transcribed from competitor reviews and industry reports, not from XPMatch's own buyers	E4, E6, E12	1	Medium	High	10 interviews per segment, recorded
3	Problem	Purchase intent for the <i>deliverable</i> is real and off-platform. First-party evidence that these buyers will pay XPMatch: none. The founders' own method classifies this as a research-slot signal, not a purchase-intent cluster	E4, E5, E12	1	Low	Existential	3 paid pilots, or 10 consistent discovery calls
4	Product	The strongest P. Working production software with the unit-cost problem solved (E2). But time-to-first-value is unmeasured, no user outside the team has completed the flow, and 8 pre-beta blockers	E1, E2, E10	2 build / 0 measured delivery	Medium-High	Medium	First-value telemetry on real accounts

#	P	Finding	Evidence	Maturity	Validation readiness	Impact	Key missing input
		are open					
5	Profit	Unit economics modelled to SDE with sensitivity. Against that: no way to collect money, no business bank account, no bookkeeping, Year-1 SDE – \$5,575, and a \$35,061 capital shortfall against the plan's requirement	E9, E11, E12	1	Low	Existential	The founders' real cash position, in writing
6	Pitch	No proof assets, no case studies, no recorded calls, no CRM, no proposal template. Two cold offer lines drafted. Evidence confirms the capability is absent	E12	0	Low	High	One delivered client to write up
7	People	Two founders with no founder agreement and no vesting, no decision-rights table, no role scorecards. Planned hires at M15 and M25 are unfunded under the shortfall	E12, [A] defaults	0	n/a	High	Executed founder agreement
8	Promote	Channels named, sequenced and given a dated kill rule — genuinely better than most pre-launch plans. Zero attributed signups, no permissioned list, no channel tested	E12	1	Medium	Medium	One 4-week channel test with referrer tracking
9	Processes	docs/OPERATING_SYSTEM.md assigns every recurring job to automation, an agent or an SOP, and the concierge SOP is written at 5 steps / 90 minutes. Never executed once	E1	1	Medium	Medium	One real concierge delivery, timed
10	Predicament	Bottleneck confirmed by data, not symptom: no payment mechanism blocks all five revenue lines at once. Continuity risk is concrete and known — one database instance, one volume, no	E10, E11	0	n/a	Existential (asset loss)	Nothing. The fix is known and overdue

#	P	Finding	Evidence	Maturity	Validation readiness	Impact	Key missing input
		backups, and unrotated keys					
11	Path	Explicit, dated and measurable: cash-positive M12, sellable M24, financial buyer as base case. Two contradictions: a year of product build preceded any selling, and the exit arithmetic is ~4× apart from this review's model	E9	2	n/a	High	Reconciliation of the ARR-multiple and SDE-multiple views
12	Protect	Nothing defined. No stated owner priorities, no intended hours, no adviser or continuity register, no decision authority if a founder is unavailable — while the plan assumes twelve unpaid months	none	? unassessed	n/a	Medium-High	Owner-stated priorities and a personal floor

Causal reading. The apparent problem is "no sales". It is not a Pitch problem. Two upstream causes produce it: **Profit** (there is no mechanism to accept money — E11) and **Problem** (no first-party evidence anyone will pay — E12). Fixing Pitch, Promote or Persona before those two would build sales machinery pointed at an unvalidated buyer with no way to transact. Equally, Product scores highest and is *not* where effort should go: more software does not move any of the three existential rows.

A cash emergency is not averaged away. Rows 5 and 10 carry existential impact and row 3 carries existential validation risk. No composite score is offered, because averaging twelve rows would hide exactly those three.

4. The two priorities

Active: Profit and Problem. Ten Ps go to the backlog in section 7 with one workflow card each.

	Why this one	What would change the choice
Profit	Existential and immediate. There is no mechanism to accept money (E11), so all five revenue lines are blocked by one missing integration. Beneath that sits a \$35,061 shortfall against the plan's own requirement — the plan is not executable as written, and no other work changes that	If the founders' real cash covers the twelve unpaid months, this drops to a two-week build task and Persona takes the slot
Problem	Existential to the thesis. Every revenue figure rests on desk research (E3–E8) and zero first-party evidence (E12). The cheapest test that can kill or confirm the whole plan lives here, and it produces the first cash	If three paid pilots close in two weeks, this converts to Pitch — turning a proven sale into a repeatable one

	Why this one	What would change the choice
	while doing it	

Why not the others. *Product* scores highest and is the wrong place to spend: more software moves none of the three existential rows. *Pitch*, *Promote* and *Persona* are all downstream of a validated problem and a working checkout — building them first aims sales machinery at an unvalidated buyer who cannot pay. *Path* has a real contradiction to resolve but it is a two-hour analysis, not a workstream, and it is listed in the 30-day review.

One dependency task, not a third workstream. docs/BETA_READINESS.md items 2 and 5 — rotate the API keys exposed during the build, and turn on Postgres backups (E10) — are folded into Priority 1 as week-0 work. One database instance with one volume and no backups means a single failure destroys the asset the entire cash plan is built on. This is protecting the subject of the workstream, not a covert extra initiative. It is hours of work, not days.

Acceptance criteria and stop conditions

Prio rity	Accountable owner	Time / budget cap	Acceptance criteria	What stops or changes it
Profi t	Tayo (unassigned finance owner for the weekly brief)	6 days of build + \$80/mo tooling	Stripe live and a real charge settled; business bank + bookkeeping open; backups on and keys rotated; a written answer on the runway gap	Founders' combined cash under \$40,000 → cut to concierge-only and keep day jobs
Prob lem	Tayo	3 weeks, \$300 (sending infrastructure)	3 concierge pilots paid or 10 discovery calls with consistent pain language and a named budget holder	Fewer than 2 pilots sold after 3 weeks of asking → willingness to pay is not there at this price; re-run against segment B only

5. Priority 1 — Profit

5.1 Artifact: 13-week cash view

Allocated from months 1–3 of 2-financial-model.md (M1 → weeks 1–4, M2 → weeks 5–8, M3 → weeks 9–13). Concierge cash lands on delivery weeks; subscription and affiliate cash is spread evenly within its month; one-time costs land in week 1. **The weekly split is an allocation assumption [A], not a forecast of any particular week.**

Wk	Concierge cash	Other cash	Total in	COGS	Opex	One-time	Net	Cash at week end
1	\$0	\$23	\$23	\$13	\$81	\$3,850	-\$3,921	\$21,079
2	\$249	\$23	\$272	\$13	\$81	—	\$178	\$21,256

Wk	Concierge cash	Other cash	Total in	COGS	Opex	One-time	Net	Cash at week end
3	\$0	\$23	\$23	\$13	\$81	—	-\$71	\$21,185
4	\$249	\$23	\$272	\$13	\$81	—	\$178	\$21,363
5	\$249	\$36	\$286	\$25	\$81	—	\$180	\$21,542
6	\$0	\$36	\$36	\$25	\$81	—	-\$70	\$21,473
7	\$249	\$36	\$286	\$25	\$81	—	\$180	\$21,652
8	\$249	\$36	\$286	\$25	\$81	—	\$180	\$21,832
9	\$249	\$64	\$313	\$34	\$65	—	\$214	\$22,046
10	\$249	\$64	\$313	\$34	\$65	—	\$214	\$22,260
11	\$0	\$64	\$64	\$34	\$65	—	-\$35	\$22,225
12	\$249	\$64	\$313	\$34	\$65	—	\$214	\$22,439
13	\$249	\$64	\$313	\$34	\$65	—	\$214	\$22,653

Lowest week-end cash: \$21,079 in week 1. Closing cash at week 13: **\$22,653**. The business does not run out of money in the first quarter — but it earns \$2,798 of cash in thirteen weeks while the founders draw nothing. That is the whole finding: this is not a cash-flow problem yet, it is a **runway** problem that arrives in month 12.

5.2 Artifact: owner-pay scenarios against the shortfall

Scenario	Owner draw	Cash needed over 12 months	Against \$25,000 available	Verdict
Plan as written	\$0 for 12 months	\$60,061 (incl. \$48,000 personal runway)	-\$35,061	Not executable
Both founders keep income elsewhere	\$0, part-time build	\$12,061 business only	+\$12,939	Executable; halve the concierge ramp and Year-1 revenue
One founder full-time, one employed	\$0 for the full-timer	~\$36,061	-\$11,061	Executable only with ~\$11k more, or a slower month 1–6
Raise the gap from a partner	\$0	\$60,061	Covered	Dilutes; contradicts the financial-buyer exit thesis

Decision required from the founders, not from this document. Scenario 2 is the only one that needs no new money, and it costs roughly half of Year-1 revenue — which, at \$40,123, is a cheaper price than it looks.

5.3 Dependency task — week 0 asset protection

Task	Source	Effort	Why it blocks everything
Turn on Postgres volume backups in Railway	E10, BETA_READINESS.md #5	Minutes	One instance, one volume, no backups: a single failure destroys the product
Rotate OpenRouter, both Google and Resend keys; redeploy after the browser key	E10, BETA_READINESS.md #2	~1 hour	Keys pasted in chat during the build are live in production
Restrict Google keys by domain and API; set spend caps	BETA_READINESS.md #3, #4	~1 hour	Worst-day exposure is ~\$45 with no cap in the app

6. Priority 2 — Problem

6.1 Artifact: pain ledger (structure, unfilled)

This table is empty on purpose. No XPMatch interview has been conducted, so filling it now would be fabrication. Every row must carry a source ID pointing at a real recording, transcript or written reply. Rows E4–E6 of the evidence ledger tell you the category transacts; they may not be copied here.

ID	Pain in the buyer's words	Segment	Source (recording / transcript ID, date)	Consequence of waiting 30 days	Current workaround	What they spend on it today	Budget holder	Verbatim or paraphrase?
P-01								
P-02								

Kill rule for the ledger: after 10 completed calls per segment, if fewer than 6 rows share the same pain *and* name a budget holder, the segment has no concentrated problem and the offer is being sold to a market that does not feel it yet.

6.2 Artifact: discovery interview guide

Ask about the last trip they planned, not about XPMatch. Never describe the product before question 8.

1. Walk me through the last trip you planned. Start from the moment you decided to go.
2. Where did you get stuck, and how long were you stuck there?
3. What did you have open when you were deciding where to eat?
4. What did you end up doing — and were you happy with it afterwards?
5. Have you ever paid anyone to help with a trip? What did you pay, and what did you get?
6. (Advisors) Walk me through building a proposal for a new client. How long, and which parts repeat?

7. (Advisors) What do you charge as a planning fee, and what do clients say about it?

8. If something produced a stay, three things to do and three restaurants matched to you in under a minute — what would make that worth paying for, and what would make it worthless?

9. What would have to be true for you to pay \$19 for one trip? \$49 for a year? \$49 a month for your practice?

10. Who else should I talk to?

Recording rule: ask permission, record, and file the transcript with an ID before the call ends. An unrecorded call produces an impression, not evidence.

6.3 Artifact: paid-test plan with decision thresholds

Test	Segment	Mechanism	Cost	Window	Pass	Kill
T1 — Concierge founding clients	A	Sell 3 trips at a founding price to the founders' own network. Money first, delivery after	\$0	Weeks 1–3	3 sold, all delivered inside 48 h	Fewer than 2 sold after 3 weeks of asking
T2 — Advisor free-first-proposal	B	400 sends/week to Fora-directory advisors; free first proposal; seat starts on the second client link	\$300	Weeks 3–9	≥5% reply and ≥2 seats started	<2% reply after list, deliverability and copy fixes
T3 — Paywall price test	A	\$19 vs \$29 Trip Pass at "Turn into a trip", 50/50, first 400 signups	\$0	Weeks 4–10	Either arm converts ≥2%	Both arms <1% → the offer, not the price, is wrong

Evidence hierarchy, enforced. A paid pilot outranks a verbal commitment with a price, which outranks a booked call, which outranks a waitlist signup, which outranks a like. Nothing below the third rung counts toward a pass.

6.4 Workflow cards

Both cards are filled in full at workflows/ and are **status: designed** — nothing has been configured, scheduled or run.

Card	P	What it does	Status
`profit-weekly-brief.yaml`	Profit	Weekly: reconcile Stripe and bank against the Year-1 monthly plan, recompute cost per active user from Google SKU counts and OpenRouter spend, raise exceptions, draft a finance brief	designed
`problem-discovery-pipeline.yaml`	Problem	On each new transcript or written reply: extract source-linked pains, urgency, payer and spend into pain-ledger rows; flag contradictions; propose follow-ups; never update the brief without passing	designed

Card	P	What it does	Status
		the quality gate	

7. Workflow backlog — the other ten Ps

One card each, not ten projects. These are **not** started; they are specified so that when a P becomes active the design already exists.

P	Workflow opportunity	Trigger	Primary output	Blocked on
Pricing	Paywall price-test tracker	Each paywall impression	Conversion by arm, by moment, with significance	Stripe live
Persona	Lead fit classifier	New advisor lead	Fit score against explicit criteria + evidence + unknowns; human decides	An ICP with first-party evidence
Product	First-value telemetry	Account events	Funnel: quiz done → package shown → swap/lock → trip created, with time-to-first-value	Activation events instrumented
Pitch	Proposal assembler	Qualified opportunity	Draft proposal from approved proof, claim-checked, routed to the sales owner	One delivered client to cite
People	Role scorecard + interview rubric builder	Role opens	Outcome scorecard, consistent questions, review packet. No automated hire/reject	A funded hiring decision
Promote	Channel attribution digest	Weekly	Signups and activation by referrer, CAC by channel, kill/scale call	Referral codes shipped
Process	Concierge delivery runner	Concierge order paid	Intake validated, draft board built from the catalog, QA checklist opened	Concierge order form
Predictment	Exception monitor	Weekly or threshold breach	Queue, concentration and cash exceptions linked to a candidate constraint	Telemetry to watch
Path	Initiative vs path check	New initiative or quarterly	KEEP / PAUSE / STOP with evidence against the documented path	— (runnable today)
Protect	Protected-time guard	Calendar request	Compare against protected blocks; draft a response; route exceptions	Owner-stated boundaries

8. Cadence

Next 7 days — owners named; dates proposed, not scheduled.

#	Action	Owner	Proposed by
1	Turn on Postgres backups; rotate all four keys; restrict Google keys and set spend caps	Tayo	Sep 24, 2026
2	Write down each founder's real available cash and pick a scenario from §5.2	Both founders	Sep 25, 2026
3	Ship Stripe billing, the three paywall moments and the concierge order form	Tayo	Oct 3, 2026
4	Ask 10 people in the founders' network to buy a concierge trip (T1)	Tayo	Oct 5, 2026
5	Open business banking and bookkeeping	unassigned [finance owner]	Oct 5, 2026
6	Buy two sending domains, set SPF/DKIM/DMARC, start warm-up (T2)	Tayo	Oct 5, 2026
7	Execute a founder agreement with vesting	Both founders	Oct 20, 2026

30-day evidence review (proposed Oct 21, 2026). Did T1 sell? What do the first pain-ledger rows say, and do they agree with docs/BUSINESS_PLAN.md §5? Is cost per active user inside \$2.00? Reconcile the ARR-multiple and SDE-multiple exit views (Path). Outcome metric: **paid pilots closed**. Countermetric: **founder hours spent per dollar collected**.

90-day reassessment (proposed Dec 21, 2026). Re-diagnose all twelve. Scale/stop: 3–5 paying clients and one repeatable channel → proceed and promote Pitch to active. Under 2% reply with no pilots → kill the advisor segment and re-run against segment C. Re-run /business-summary in Update mode and compare actuals to the Year-1 monthly plan.

These dates are text, not reminders. No scheduler has been configured. Say the word and I will set a recurring check-in with the tooling available in this session.

9. Execution receipt

What	State	Evidence
Twelve-P diagnosis	Done — provisional, from p-map.md	§3 above
Business Brief and evidence ledger	Done	§1, §2
13-week cash view	Done, derived from the verified model	§5.1; reconciles to 2-financial-model.xlsx
Owner-pay scenarios	Done	§5.2
Pain ledger	Structure only, deliberately empty	§6.1 — no interview exists to fill it

What	State	Evidence
Interview guide, paid-test plan	Done	§6.2, §6.3
Two workflow cards	Designed. Not drafted, dry-run tested, configured or live	workflows/*.yaml
Trial on supplied data	Blocked. Exact missing input: one recorded discovery call or one settled Stripe charge	—
Backlog for the other ten Ps	Done	§7
Scheduled reminders	Not created. Dates above are proposals	—

What did not run, and why. The Problem pipeline cannot be trialled: there is no transcript to feed it, and generating a synthetic one would put invented customer language into a ledger meant to hold only real quotes. The Profit brief cannot be trialled: there is no Stripe account, no bank feed and no accounting export. Both become runnable the week Stripe goes live and the first call is recorded.

Remaining access needs: Stripe (read), business bank (read), Google Cloud billing (read), OpenRouter usage (read), a transcript store, and a scheduler. None is connected in this session.

To resume: re-run own - run after the 30-day review with the first pain-ledger rows and the first Stripe settlement attached, and it will re-diagnose against real evidence instead of defaults.